

DDPM

marii

Add explanation of why subclassing `TensorImageBase`, and what is `TypeDispatch`. (title needs to be full name)

We create two new Tensor Types, one for our noise, and one for our timestep.

```
class TensorNoise(TensorImageBase):pass
class TensorStep(TensorBase): pass
```

We would like normalize to denormalize our noise before showing it. This is so the noise in our image looks similar to the noise in our noised image.

```
@Normalize
def decodes(self, x:TensorNoise):
    f = to_cpu if x.device.type=='cpu' else noop
    return (x*f(self.std) + f(self.mean))

norm = Normalize.from_stats(*imagenet_stats)
show_images(norm.decode(img))
```



I patch `ItemTransform` here, so that it can work off of `TypedTuples`. Essentially if we have a `DiffusionTuple`, the transform will apply to that if it should apply to that type of tuple.

```

class ItemTransform(Transform):
    "A transform that always take tuples as items"
    _retain = True
    # Only showing important code
    def _call1(self:ItemTransform, x, name, **kwargs):
        if not _is_tuple(x): return getattr(super(), name)(x, **kwargs)
        y=self._call_tuple(name,x,**kwargs)
        if not self._retain: return y
        if is_listy(y) and not isinstance(y, tuple): y = tuple(y)
        return retain_type(y, x)
    def _call_tuple(self:ItemTransform, name, x, split_idx=None, **kwargs):
        f = getattr(super(), name)
        f2name='encodes' if name == '__call__' else 'decodes' if name == 'decode' else name
        f2 = getattr(self, f2name)
        if isinstance(f2,TypeDispatch) and f2[type(x)] is not None:
            if split_idx!=self.split_idx and self.split_idx is not None: return x
            y = f2(x, **kwargs)
        else:
            y = f(list(x), **kwargs)
        return y

```

The general idea is to implement a named tuple, and use duck typing. In the future, we should look at the named tuple class and do something more similar to that.

```

class DiffusionTuple(fastuple):
    def __new__(cls, *rest):
        self=super().__new__(cls, *rest)
        i=0
        self.x=self[i]
        if(isinstance(self[i+1],TensorImage)): self.x0=self[i:=i+1]
        self.t=self[i:=i+1]
        if(len(self)>i+1): self.y=self[i:=i+1]
        if(len(self)>i+1): self.pred=self[i:=i+1]
        if(len(self)>i+1): self.sampled_pred=self[i:=i+1]
        return self

```

A little transform to make our tuple a DiffusionTuple

```

class ToDiffusionTuple(ItemTransform):
    order=100
    def encodes(self,xy):

```

```
return DiffusionTuple(*xy[:-1],TensorNoise(xy[-1]))
```

This `Transform` expects `y` to contain an image, and just replaces it with noise. Our model tries to predict the noise in an image.

```
class LabelToNoise(ItemTransform):
    order=101
    def encodes(self,xy:DiffusionTuple):
        y=xy.y
        xy.y[:]=TensorNoise(torch.randn_like(y))
        return xy
```

```
diff_tuple=LabelToNoise.encodes(DiffusionTuple(img[0].detach().clone(),TensorStep(torch.te
```

We can access tuple elements by attributes. This is useful when you don't know what index a particular value is located.


```
diff_tuple.x.shape,diff_tuple.t.shape,diff_tuple.y.shape
```

```
(torch.Size([3, 320, 480]), torch.Size([1, 1]), torch.Size([3, 320, 480]))
```

Now we have a way to create an image, and convert the label to noise.

```
norm.decode(diff_tuple).show(show_noise=True)
```

```
<AxesSubplot:title={'center': 'TensorStep([[200]])'}>
```



Next, we need to go create a noised image, to pass to our model.

First, how much noise to apply to each step?

```

class LinearNoiseSchedule:
    "Schedule like used in DDPM"
    def __init__(self, betas=None, n_steps=None, device='cuda'):
        if betas is not None: self.n_steps=betas.shape[0]
        if n_steps is None: self.n_steps=1000
        if betas is None: self.betas = torch.linspace(0.0001, 0.02, self.n_steps, device=device)
        self.alphas = 1. - self.betas
        self.alpha_bar = torch.cumprod(self.alphas, dim=0)

```


Lets graph the various values here, in order to see what happens. Pay particularly close attention to `alpha_bar` as that controls the balance between our signal(image) and our noise.

```

lns=LinearNoiseSchedule()
plt.plot((lns.betas).cpu())
plt.plot((lns.alphas).cpu())
plt.plot((lns.alpha_bar).cpu())
plt.legend(['betas', 'alphas', 'alpha_bar'])

```

<matplotlib.legend.Legend at 0x7f89721ee710>



Next is DDPM-style Q-sampling. This is pretty much used for all diffusion models, and is the process that takes us from an image to noise.

```
class DDPM_Q_Sampling():
    def __init__(self, predicts_x=False, noise_schedule=LinearNoiseSchedule(), n_steps=1000, device=device):
        self.device=device
        self.ns=noise_schedule
        self.n_steps=n_steps
        self.t_sched=torch.linspace(0, len(self.ns.alpha_bar)-1, n_steps, dtype=torch.long)
    def __call__(self, x, es, t):
        t=self.t_sched[t]
        a=self.ns.alpha_bar[t].to(device=x.device)
        signal = (a ** .5)*x
        noise = (1-a)**.5 * es
        return signal + noise
    def undo(self, z, es, t):
        "Goes back to the original image given noise. Only works if es is the original noise"
        if(isinstance(es, TensorImage)):
            return es
```

```

t=self.t_sched[t]
a=TensorBase(self.ns.alpha_bar[t].to(device=z.device))
noise=TensorBase((1-a)**.5 * es)
return TensorImage((z-noise)/(a **.5))

```

```
diff_trans = DiffusionSamplingTransform(DDPM_Q_Sampling(),lambda x:x)
```

```
norm.decode(diff_trans(diff_tuple)).show()
```

```
<AxesSubplot:title={'center':'TensorStep([[200]])'}>
```




Lets now test so make sure our noise is being generated correctly.

```
noise_tuple=LabelToNoise.encodes(DiffusionTuple(img[0].detach().clone(),TensorStep(torch.t
```

```
norm.decode(diff_trans(noise_tuple)).show(show_noise=True)
```

```
<AxesSubplot:title={'center':'TensorStep([[999]])'}>
```

TensorStep([[999]])



These are not exactly the same as it is one noising step, but they are fairly close.

```
is_close(norm.decode(diff_trans(noise_tuple))[0],TensorImage(norm.decode(diff_trans(noise_
```

```
TensorImage(True, device='cuda:0')
```

Going from noise to and image, p_sampling

P sampling is going from noise to image!

```
class Diffusion_P_Sampler():
    def __init__(self,model,sampling_function):
        self.device=sampling_function.device
        self.model=model
        self.sampling_function=sampling_function
    # __call__ implemented, but not shown.
    def iter_noise(self,x_t,ts,t_start):
        i=0
        while((ts>0).any()):
            x,t=x_t[ts>0],ts[ts>0]
            with autocast(device_type=self.device, dtype=x.dtype):
                with torch.no_grad():
                    e = self.model(x,self.deconvert(t) if i!=0 else t_start)
                    x_t[ts>0]=self.sampling_function(x,e,t,t=t_start if i==0 else None)
            ts[ts>0]-=1
            i+=1
        yield x_t
```

Notice here that we generate noise as random numbers during our sampling process. This makes DDPM sampling a stochastic process.

```

@patch
def __call__(self:DDPM_P_Sampling,x,es,ns_t,t=None):
    t= self.t_sched[ns_t] if(t is None) else t[... ,None,None,None]
    n=torch.randn_like(x)
    e,a,b=self._noise_at_t(es,t),self.ns.alphas[t],self.ns.betas[t]
    signal = (x - e) / (a ** 0.5)
    noise = b**.5 * n
    return signal + noise

@patch
def _noise_at_t(self:DDPM_P_Sampling,es,t):
    eps_coef = (1 - self.ns.alphas[t]) / (1 - self.ns.alpha_bar[t]) ** .5
    return eps_coef* es

```

We implement DDIM sampling here, as it drastically reduces sampling time from 1000 steps to 50. DDIM is also a deterministic sampler, meaning we do not have random numbers generated as part of our sampling process. Just generally helps us keep our sanity when trying to show our results. <https://arxiv.org/abs/2010.02502>

```

@patch
def __call__(self:DDIM_P_Sampling,z,es,ns_t,t=None):
    if(t is None): t=self.t_sched[ns_t]
    tp1=self.t_sched[ns_t-1]
    a,a_tp1=self.ns.alpha_bar[t][... ,None,None,None],self.ns.alpha_bar[tp1][... ,None,None,]
    if self.predicts_x:
        xs=es
        es=(z - (a)**.5 * xs)/(1-a)**.5
    else: xs=(z - (1-a)**.5 * es)/ (a ** .5)
    signal = a_tp1**.5*(xs)
    noise = (1-a_tp1)**.5*es
    return signal + noise

```

Training a model

```

path = untar_data(URLs.CIFAR)

m=Unet(dim=192+192//8,channels=3,).cuda()

bs=128
n_steps=1000

```



```

diffusion_transform = DiffusionSamplingTransform(DDPM_Q_Sampling(),Diffusion_P_Sampler(m,D
dls=DataBlock((ImageBlock(),
                TransformBlock(type_tfms=[Transform(enc=lambda o: TensorStep(o))]),
                ImageBlock()),
                n_inp=2,
                item_tfms=[Resize(32)],
                batch_tfms=(Normalize.from_stats(*cifar_stats),ToDiffusionTuple,LabelToNoise,dif
                get_items=get_image_files,
                get_x=[lambda x:x,
                        lambda x: torch.randint(1, n_steps, (1,), dtype=torch.long)],
                splitter=IndexSplitter(range(bs)),
                ).dataloaders(path,bs=bs,val_bs=2*bs)
dls.show_batch()







```

```

/home/molly/miniconda3/envs/fastai/lib/python3.10/site-packages/torch/_tensor.py:1121: UserW
ret = func(*args, **kwargs)

```

Input/Original/Target



```
def mse_loss_weighted(ys,targ):
    return torch.mean(targ.w_sched[...,None] * ((ys - targ).flatten(start_dim=1) ** 2))


def snr(at): return at/(1-at)
```

Taken from a paper called “Perception Prioritized Training of Diffusion Models, this is a continuous version of the weights. Which becomes our signal=to-noise ration, over the change in our signal to noise ratio.”<https://arxiv.org/abs/2204.00227>

```
def continuous_weights(at):
    weights = -snr(at[1:])/((snr(at[1:])-snr(at[:-1])))
    return torch.cat((weights[0:1],weights))
```

This is what the weights look like. Notice I clip the weights at 1.

```
plt.plot(continuous_weights(LinearNoiseSchedule().alpha_bar).cpu().clip(min=1))
```



```

class WeightedLinSched(Callback):
    def after_pred(self):
        if(not hasattr(self,'ws')):
            self.ws = continuous_weights(LinearNoiseSchedule().alpha_bar).clip(min=1)
            self.ws /= self.ws.mean()
        ts=self.learn.xb[1].flatten()
        self.learn.yb[0].w_sched=self.ws[ts]

learn = Learner(dls,m,mse_loss_weighted,opt_func=Adam,cbs=[FlattenCallback,WeightedLinSched])
learn = learn.to_fp16()
learn.fit_flat_cos(10,lr=2e-4,wd=1e-4)

```

<IPython.core.display.HTML object>

epoch	train_loss	valid_loss	time
0	0.049204	0.042475	04:47
1	0.043398	0.044983	04:48
2	0.042191	0.034068	04:49
3	0.040477	0.038086	04:48
4	0.039220	0.038858	04:49
5	0.038735	0.032376	04:49
6	0.038864	0.029669	04:49
7	0.038684	0.041698	04:49
8	0.037839	0.031916	04:50
9	0.037525	0.041718	04:50

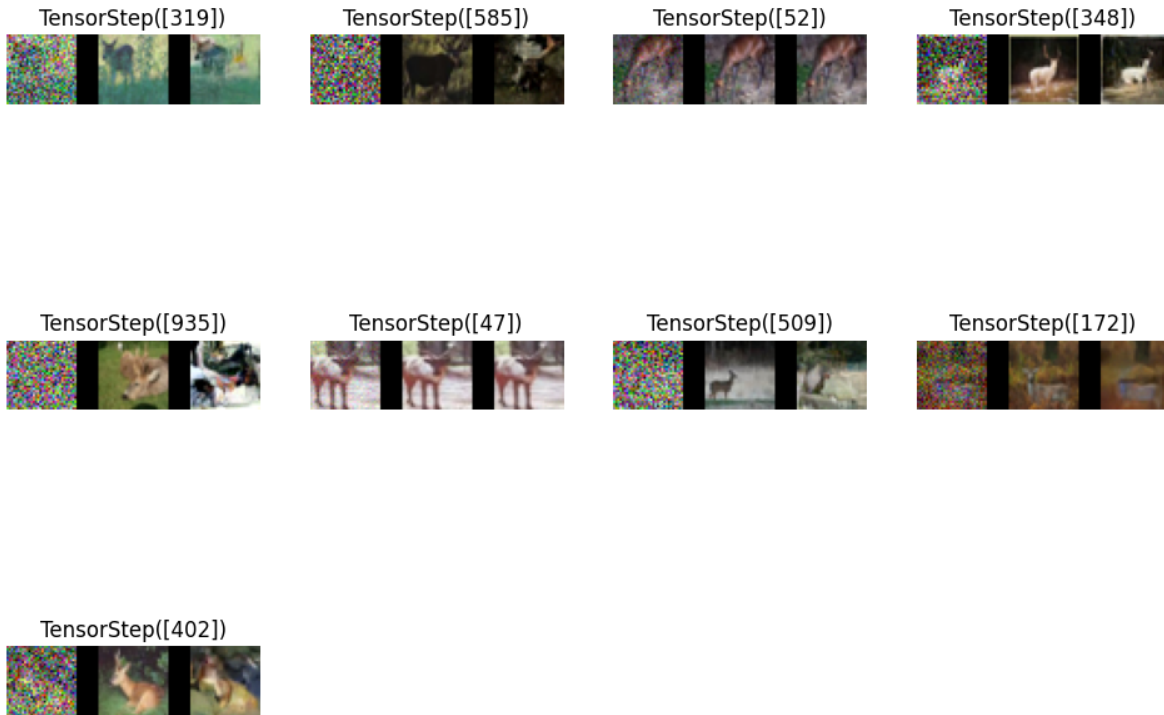
next check show_results

```
learn.show_results()
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Input/Original/DenoisedImage




```
learn.show_results(show_noise=True)
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Input/Original/DenoisedImage



channels last, fused optimizers(foreach benjamin's), jit model.